

Reinforcement Learning in Healthcare:

Designing an Intelligent Insulin Dosing System

Marvin Azuogu

Houston City College

ITAI 4375- AI in Healthcare-RT-17185

Spring 2026

Professor: Dr. Doreen Rosenstrauch

March 30, 2026.

What Is Reinforcement Learning?

Reinforcement Learning (RL) is a branch of artificial intelligence (AI) in which a computer program, called an agent, learns to make better decisions over time by interacting with its environment. Rather than following a fixed set of pre-programmed rules, the agent experiments with different actions and receives feedback in the form of rewards or penalties. When an action leads to a good outcome, the agent receives a positive reward; when it leads to a poor outcome, it receives a penalty. By repeating this process many times, the agent gradually learns which actions produce the best results in which situations (Lee et al., 2024).

A simple analogy is training a dog: when the dog performs the correct behavior, it earns a treat; when it misbehaves, it does not. Over time, the dog learns to repeat rewarded behaviors. An RL agent works the same way, except its environment can be as complex as a hospital ward and its "treat" is a mathematical reward signal. Because RL systems continuously improve with experience, they are especially well suited to healthcare, where conditions change moment-to-moment and the optimal treatment decision may differ from patient to patient (Alshammari et al., 2024).

Healthcare Scenario: Adaptive Insulin Dosing for Type 1 Diabetes

The scenario chosen for this report is the automated management of insulin delivery for patients with Type 1 Diabetes Mellitus (T1DM) admitted to a hospital intensive care unit (ICU). Patients with T1DM cannot produce their own insulin, so their blood glucose must be controlled around the clock. Too little insulin results in dangerously high blood sugar (hyperglycemia), while too

many causes a life-threatening drop in blood sugar (hypoglycemia). In current practice, nurses or physicians adjust insulin infusion rates based on periodic glucose checks, but human schedules and workload make it difficult to achieve continuous, personalized control.

An RL agent connected to a continuous glucose monitor (CGM) and an automated insulin pump could observe the patient's glucose level every five minutes and decide in real time how much insulin to administer, adapting instantly to meals, physical activity, stress, or changes in the patient's condition. This concept is sometimes called an "artificial pancreas" and represents one of the most clinically impactful near-term applications of RL in medicine (Lee et al., 2024).

Components of the RL System

A. State

The state is all the information the RL agent uses to understand the patient's current situation at each decision point. In this scenario, the state would include:

- Current blood glucose level (mg/dL), measured by the CGM every five minutes
- Rate of change in blood glucose over the past 15–30 minutes (rising, stable, or falling)
- Total insulin delivered in the past hour, to prevent accidental overdosing
- Recent carbohydrate intake as logged by the nursing staff
- Time of day, since insulin sensitivity follows a circadian rhythm
- Patient-specific factors such as age, weight, and history of hypoglycemic episodes

Together, these data points give the agent a rich, real-time picture of the patient's metabolic state, far more detail than a nurse reviewing a chart at the top of the hour.

B. Action

The action is the decision the RL agent makes at each time step. For this insulin dosing system, the available actions would be:

- Increase the insulin infusion rate by a small, medium, or large increment
- Decrease the insulin infusion rate by a small, medium, or large increment
- Maintain the current infusion rate unchanged
- Suspend insulin delivery temporarily if severe hypoglycemia is detected or predicted

All actions would be constrained within clinically safe boundaries defined by physicians.

Regardless of what the agent's learning suggests, it could never exceed a maximum safe dose, human defined safety rules always take priority.

C. Reward

The reward function defines what "success" means for the agent. For the insulin dosing system, reward would be structured around keeping the patient's blood glucose in the clinical target range of 70–180 mg/dL:

- Positive reward (+1 to +5): Blood glucose is within the target range; the closer to the ideal center of approximately 100–140 mg/dL, the higher the reward
- Mild penalty (−1): Blood glucose is slightly outside the target range (mild hyperglycemia or mild hypoglycemia)
- Severe penalty (−10 or more): Blood glucose falls below 54 mg/dL (severe hypoglycemia) or rises above 300 mg/dL (severe hyperglycemia), as both are potentially fatal

This reward structure teaches the agent that small fluctuations are tolerable, moderate deviations are undesirable, and extreme glucose values must be avoided at almost any cost, closely mirroring the priorities of real clinical care (Alshammari et al., 2024).

How the RL System Could Improve Patient Care

Unlike fixed sliding-scale protocols that apply the same rules to every patient, an RL-based system would personalize care over time. As it interacts with a specific patient over hours and days, it would learn that individual's unique insulin sensitivity, typical glucose patterns, and responses to meals. It would also respond far faster than human staff: a nurse might check glucose every one to four hours, whereas the RL agent acts every five minutes, catching dangerous trends before they escalate into emergencies.

Research supports the promise of this approach. Lee et al. (2024) published a clinical primer in *npj Digital Medicine* confirming that RL algorithms have demonstrated meaningful improvements in glycemic control, noting that RL-based controllers for type 2 diabetes achieved

statistically significant reductions in HbA1c compared to standard care in proof-of-concept trials. Alshammari et al. (2024) similarly noted in their MDPI Sensors review that RL techniques can dynamically adapt to the complexity of real-world healthcare environments, reducing patient harm events and improving outcomes across multiple ICU applications. With tighter glucose control, patients may experience shorter ICU stays, fewer complications such as infection or organ damage, and reduced overall healthcare costs.

Potential Challenge: Safety During Learning

One of the most significant concerns in applying RL to insulin dosing, or any life-critical clinical task, is ensuring patient safety during the agent's learning process. RL agents improve by exploring different actions, including ones that may turn out to be suboptimal. In a video game, a wrong move simply restarts the level. In an ICU, a wrong insulin dose could trigger a seizure, cardiac arrhythmia, or death.

This is why RL systems for healthcare must be trained extensively in validated computer simulations, such as the FDA-accepted UVA/Padova metabolic simulator for diabetes research, before being tested on real patients. Even after deployment, the system must operate within strict physician-defined safety boundaries and always defer to clinician override. There is also the challenge of explainability: nurses and doctors need to understand why the system chose a particular dose. Lee et al. (2024) explicitly identify interpretability as one of the four central barriers to RL deployment in medicine, noting that "black box" decision processes undermine clinician trust and complicate regulatory approval. Addressing this requires close collaboration among AI researchers, clinicians, ethicists, and regulators.

Conclusion

Reinforcement Learning holds extraordinary potential to transform hospital care. By personalizing decisions to each individual patient and responding continuously in near real time, RL systems can provide a level of vigilance that would be impossible for any single human clinician to replicate consistently. The insulin dosing scenario described in this report illustrates how the three core components, state, action, and reward, can be designed to align a machine's learning objectives with the priorities of real clinical medicine.

At the same time, the challenges must not be underestimated. Safety, transparency, and clinician trust are not optional, they are prerequisites for any RL system influencing life-and-death decisions. As simulation environments improve, regulatory pathways for AI-driven clinical tools mature, and interpretability methods advance, RL-powered decision support is likely to become a standard part of hospital care within the coming decade. Used responsibly, this technology has the potential to save lives, reduce the burden of chronic disease, and help healthcare providers deliver better outcomes for every patient they serve.

References

Alshammari, H., Azzeh, M., Al-Khatib, R., Alzaatreh, A., & Batayneh, W. (2024).

Reinforcement learning algorithms and applications in healthcare and robotics: A comprehensive and systematic review. *Sensors*, 24(8), 2461.

<https://doi.org/10.3390/s24082461>

Lee, J., Jin, W., Kim, H., & Park, S. (2024). A primer on reinforcement learning in medicine for clinicians. *npj Digital Medicine*, 7, 315. <https://doi.org/10.1038/s41746-024-01316-0>