Evaluating Federated Learning Case Studies in Healthcare

Marvin Azuogu

Houston City College

ITAI 4375- AI in Healthcare-RT-17185

Spring 2026

Professor: Dr. Doreen Rosenstrauch

April 27, 2026.

Introduction

Federated learning is a machine learning technique that enables multiple healthcare institutions to collaboratively train AI models without sharing sensitive raw patient data. Instead of centralizing all data in one location, federated learning keeps data locally at each institution while sharing only model parameters and insights across a network. This approach is critically important in healthcare because it allows researchers to build more robust, generalizable models using diverse patient populations while maintaining strict data privacy and compliance with regulations like HIPAA (Health Insurance Portability and Accountability Act). This report examines two landmark case studies that demonstrate how federated learning is transforming healthcare research: one focused on privacy-first health research methodologies, and another on predicting COVID-19 mortality outcomes.

Case Study 1: Privacy-First Health Research with Federated Learning

Overview and Main Goal

This study, conducted by Sadilek et al. (2021) and published in NPJ Digital Medicine, aimed to demonstrate that federated learning combined with differential privacy could enable large-scale health research studies across diverse populations while maintaining superior privacy protections compared to traditional centralized approaches. The primary objective was to prove that federated models could achieve comparable accuracy, precision, generalizability, and interpretation as standard centralized statistical models without significantly increasing computational costs.

Healthcare Data Used

The study utilized diverse health data types collected across multiple sites and participants, including clinical and epidemiological data from individual health records. The researchers applied federated learning to single and multi-site health studies, testing the methodology on various data modalities including patient demographics, medical histories, health events, and real-time health signals from individuals' devices. This diverse dataset reflected real-world healthcare scenarios where patient data is naturally distributed across different healthcare systems and institutions (Sadilek et al., 2021).

Key Findings and Results

The study demonstrated several groundbreaking findings. First, federated models achieved similar accuracy, precision, and generalizability compared to standard centralized statistical models, proving that privacy protection does not require sacrificing model performance. Second, the researchers successfully incorporated differential privacy, a mathematical technique that adds controlled noise to data, directly into the federated learning approach, providing considerably stronger privacy protections. Third, the federated models achieved these results without significantly raising computational costs compared to centralized approaches. Importantly, this was the first work to apply modern federated learning methods explicitly incorporating differential privacy to clinical and epidemiological research across diverse health studies (Sadilek et al., 2021). This groundbreaking approach demonstrated that privacy and accuracy are not mutually exclusive goals in healthcare AI.

Case Study 2: Federated Learning of Electronic Health Records to Improve Mortality Prediction in COVID-19 Patients

Overview and Main Goal

This research, conducted by Vaid et al. (2021) and published in JMIR Medical Informatics, was the first study to build federated learning models to predict mortality in hospitalized COVID-19 patients within seven days of hospital admission. The goal was to develop robust predictive models that could provide early warning of high-risk patients across multiple healthcare institutions without requiring the sharing of raw patient data, which is particularly important given the sensitive nature of pandemic research.

Healthcare Data Used

The study utilized electronic health record (EHR) data from 4,029 COVID-19 patients across five hospitals within the Mount Sinai Health System in New York City. The EHR data included patient demographics, past medical history, admission vitals (such as blood pressure, heart rate, and temperature), and laboratory values (such as blood tests and oxygen levels). These data were diverse across the five hospitals, reflecting real-world variations in patient populations, demographics, outcomes, sample sizes, and lab values. Two machine learning models were trained: logistic regression with L1 regularization (LASSO) and multilayer perception (MLP) neural networks. Models were compared in three configurations: local models (using only individual hospital data), pooled models (combining data from all five sites), and federated models (training distributed data without combining it) (Vaid et al., 2021).

Key Findings and Results

The study produced several important findings. Both the federated LASSO and federated MLP models performed better than their local model counterparts at four of the five hospitals, demonstrating the benefit of collaborative learning across institutions. The federated MLP model outperformed the federated LASSO model at all hospitals, showing that neural networks captured complex relationships in the data more effectively. Critically, the federated models achieved performance comparable to pooled models (which did combine raw data) without any hospital having to share sensitive patient information. The study demonstrated that federated learning allows hospitals with diverse patient populations and characteristics to benefit from collaborative research while maintaining strict data privacy. The researchers identified several important risk factors for COVID-19 mortality, including past medical history (such as pneumonia, diabetes, and cancer), and showed that federated learning successfully identified these patterns across multiple institutions (Vaid et al., 2021).

Comparison of the Two Case Studies

Similarities

Both case studies share several important characteristics. Both explicitly prioritize patient privacy as a core research objective, implementing federated learning to prevent raw data from leaving individual healthcare institutions. Both studies demonstrate that federated models can achieve comparable or better performance than centralized approaches. Both research projects were conducted in 2021, representing a pivotal moment when federated learning was emerging as a practical solution for healthcare research. Additionally, both studies involved multi-site healthcare collaborations, Sadilek et al. worked across diverse health systems and populations, while Vaid et al. collaborated across five hospitals in one health system. Finally, both studies emphasize the real-world applicability of federated learning, showing that it works with actual healthcare data, not just simulated datasets.

Differences

The two studies differ in important ways. Sadilek et al. (2021) focused on establishing the general methodological framework, proving that federated learning with differential privacy could work across diverse health studies with various data types and research questions. This was foundational research demonstrating the viability of the approach. In contrast, Vaid et al. (2021) applied federated learning to a specific, urgent clinical problem: predicting mortality in COVID-19 patients within a narrow time window (seven days). Sadilek et al. worked with diverse health data types across multiple research studies, while Vaid et al. focused specifically on EHR data from a single health system (though across multiple hospitals). The scope also differed, Sadilek et al. demonstrated general applicability across many different health conditions and study types, while Vaid et al. demonstrated practical clinical utility for a specific infectious disease prediction task. Additionally, Vaid et al. explicitly compared federated models against pooled models, demonstrating that federated approaches matched the performance of models trained on centralized data.

Main Benefits of Federated Learning in Each Study

For Sadilek et al. (2021), the main benefits included: (1) Stronger privacy protections through differential privacy without sacrificing model accuracy, (2) Scalability, enabling research studies that can include large population numbers without centralized data storage, (3) Flexibility, demonstrating that the approach works across diverse health datasets and research questions, and (4) Participant confidence, individuals could contribute to health research knowing their raw data never leaves their device or healthcare system. For Vaid et al. (2021), the key benefits included: (1) Improved model performance through access to diverse patient data across multiple hospitals without exposing sensitive information, (2) Practical clinical utility, enabling hospitals to benefit from collaborative research during a pandemic emergency, (3) Regulatory compliance, avoiding the complexity of centralized data aggregation and associated HIPAA requirements, and (4) Rapid deployment, allowing models to be trained and shared across hospitals quickly without lengthy data-sharing agreements.

Challenges and Limitations of Federated Learning

Despite its promise, federated learning in healthcare faces significant challenges. One major limitation evident from these studies is data heterogeneity, different healthcare institutions collect data differently, use different medical coding systems, measure variables in different ways, and serve patients with different demographics. This heterogeneity makes it difficult to build models that perform well across all sites. Additionally, while federated learning protects raw data, researchers have discovered that model parameters and updates shared during federated learning can potentially leak information about institutional data through sophisticated attacks. Implementing additional privacy protections like differential privacy (as Sadilek et al. did) adds computational overhead and can reduce model accuracy if not carefully calibrated. Furthermore, federated learning requires robust infrastructure, technical expertise, and governance frameworks that many healthcare institutions lack. Organizations must establish protocols for model training, validation, communication between sites, and handling of model errors, all while maintaining security and privacy. Computational costs, though not prohibitively high according to Sadilek et al., can still be substantial for real-time applications. Finally, regulatory uncertainty remains, healthcare regulators are still developing guidelines for federated learning validation, model transparency, and accountability when models trained across multiple sites produce errors affecting patient care.

Future Impact of Federated Learning on Healthcare Research

Federated learning has the potential to fundamentally transform healthcare research. These two case studies suggest several transformative impacts. First, federated learning will likely accelerate medical discoveries by enabling researchers to build better models using larger, more diverse patient populations without the barriers of centralized data governance. Second, it democratizes access to AI research, smaller healthcare institutions that previously couldn't participate in large studies due to data privacy concerns can now contribute. Third, federated learning addresses a persistent privacy paradox in healthcare: the tension between wanting to use data to help more patients and protecting individual privacy. By decoupling data sharing from research collaboration, federated learning resolves this tension. Fourth, it can accelerate responses to public health emergencies (as demonstrated by the COVID-19 mortality prediction work) by enabling rapid model training across multiple sites without delays from data-sharing agreements. Looking forward, we can expect federated learning to become standard practice in healthcare, particularly as technical barriers decrease, regulatory frameworks clarify, and healthcare institutions invest in the infrastructure required. Emerging approaches combining federated learning with other privacy-enhancing technologies (like homomorphic encryption and secure multi-party computation) may provide even stronger protection. Additionally, federated learning could extend beyond traditional research to support clinical decision-making tools, models trained on federated data could be deployed to individual hospitals and continuously improved as new local data emerges. The combination of federated learning with explainability techniques could also help clinicians understand and trust AI recommendations, addressing a critical barrier to clinical adoption. Ultimately, federated learning represents a paradigm shift from asking institutions to surrender data for greater good to asking them to collaborate while maintaining control over their data, a more sustainable and ethically sound approach to healthcare AI research.

Conclusion

These two federated learning case studies represent watershed moments in healthcare AI. Sadilek et al. (2021) proved that federated learning combined with differential privacy could work in practice, achieving privacy and accuracy without unacceptable computational cost. Vaid et al. (2021) demonstrated that federated learning could solve an urgent real-world clinical problem, predicting COVID-19 mortality, better than isolated institutional efforts. Together, they establish federated learning as more than a theoretical concept; it is a practical, deployable technology that addresses one of healthcare's most pressing challenges: how to harness the power of AI and big data while protecting patient privacy and maintaining public trust. The path forward is clear but challenging. Healthcare organizations must invest in technical infrastructure, develop governance frameworks, and establish regulatory clarity. Researchers must continue testing federated approaches in diverse clinical settings and identifying and solving technical limitations. Regulators must develop frameworks that encourage innovation while ensuring patient safety and privacy. Patients must be informed about and involved in decisions about how federated learning will be used in their care. The potential benefits are enormous, better models, faster research, democratized participation, and a privacy-protecting approach to healthcare AI. But realizing these benefits requires commitment from all stakeholders to overcome current limitations and challenges. If successful, federated learning could become as transformative to healthcare research in the coming decade as centralized databases were in the previous one.

References

Sadilek, A., Liu, L., Nguyen, D., Kamruzzaman, M., Serghiou, S., Rader, B., Ingerman, A., Mellem, S., Kairouz, P., Nsoesie, E. O., MacFarlane, J., Vullikanti, A., Marathe, M., Eastham, P., Brownstein, J. S., Aguera y Arcas, B., Howell, M. D., & Hernandez, J. (2021). Privacy-first health research with federated learning. NPJ Digital Medicine, 4, 132. https://doi.org/10.1038/s41746-021-00489-2

Vaid, A., Jaladanki, S. K., Xu, J., Teng, S., Kumar, A., Lee, S., Somani, S., Paranjpe, I., De Freitas, J. K., Wanyan, T., Johnson, K. W., Bicak, M., Klang, E., Kwon, Y. J., Costa, A., Zhao, S., Miotto, R., Charney, A. W., Böttinger, E., ... Glicksberg, B. S. (2021).

Federated learning of electronic health records to improve mortality prediction in hospitalized patients with COVID-19: Machine learning approach. JMIR Medical Informatics, 9(1), e24207. https://doi.org/10.2196/24207

Kaissis, G. A., Makowski, M. R., Rückert, D., & Braren, R. F. (2020). Secure, privacy-preserving and federated machine learning in medical imaging. Nature Machine Intelligence, 2(6), 305–311. https://doi.org/10.1038/s42256-020-0186-1

Sheller, M. J., Reina, G. A., Edwards, B., Martin, J., & Bakas, S. (2020). Federated learning in medicine: Facilitating multi-institutional collaborations without sharing patient data. Scientific Reports, 10, 12598. https://doi.org/10.1038/s41598-020-69250-1

Yang, Q., Liu, Y., Chen, T., & Tong, Y. (2024). Federated learning: Challenges, methods, and future directions. IEEE Transactions on Information Forensics and Security, 19, 4833–4847. https://doi.org/10.1109/TIFS.2023.3345594